


Data LIBERATION and  
data INTEGRATION


with  Apache Kafka

---

Martin Kleppmann @martinkl



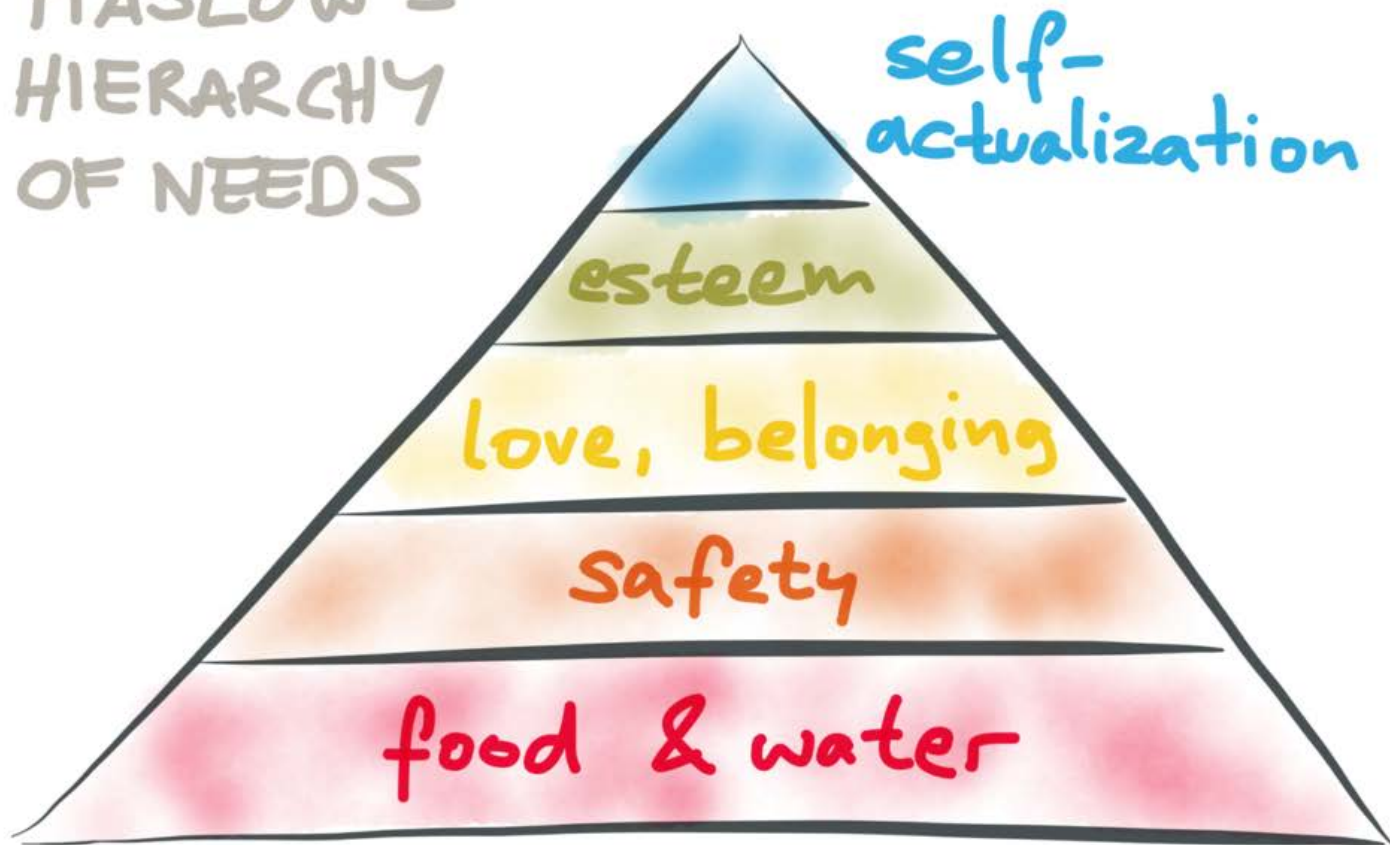
UNIVERSITY OF  
CAMBRIDGE




dataintensive.net

@martinkl


# MASLOW'S HIERARCHY OF NEEDS



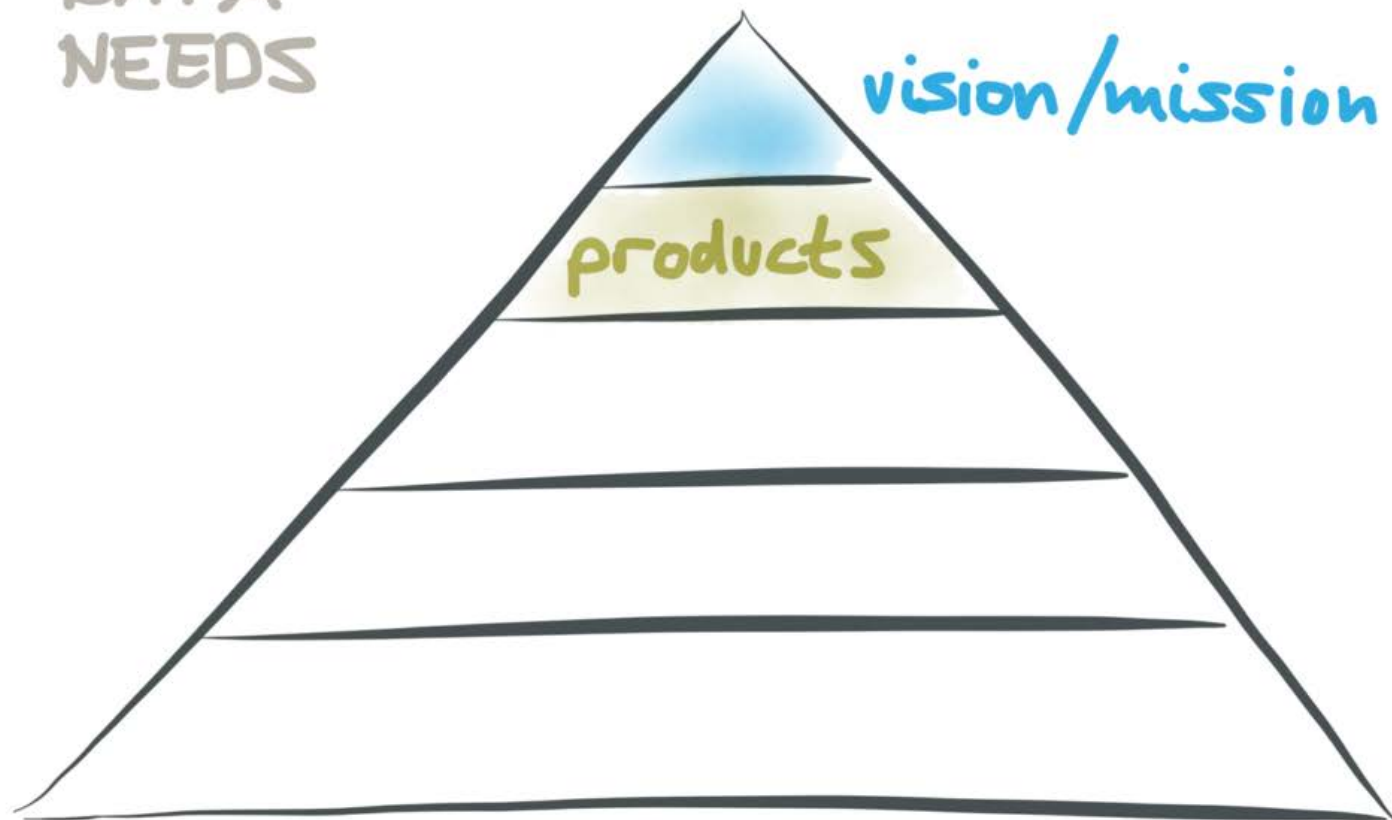
DATA  
NEEDS



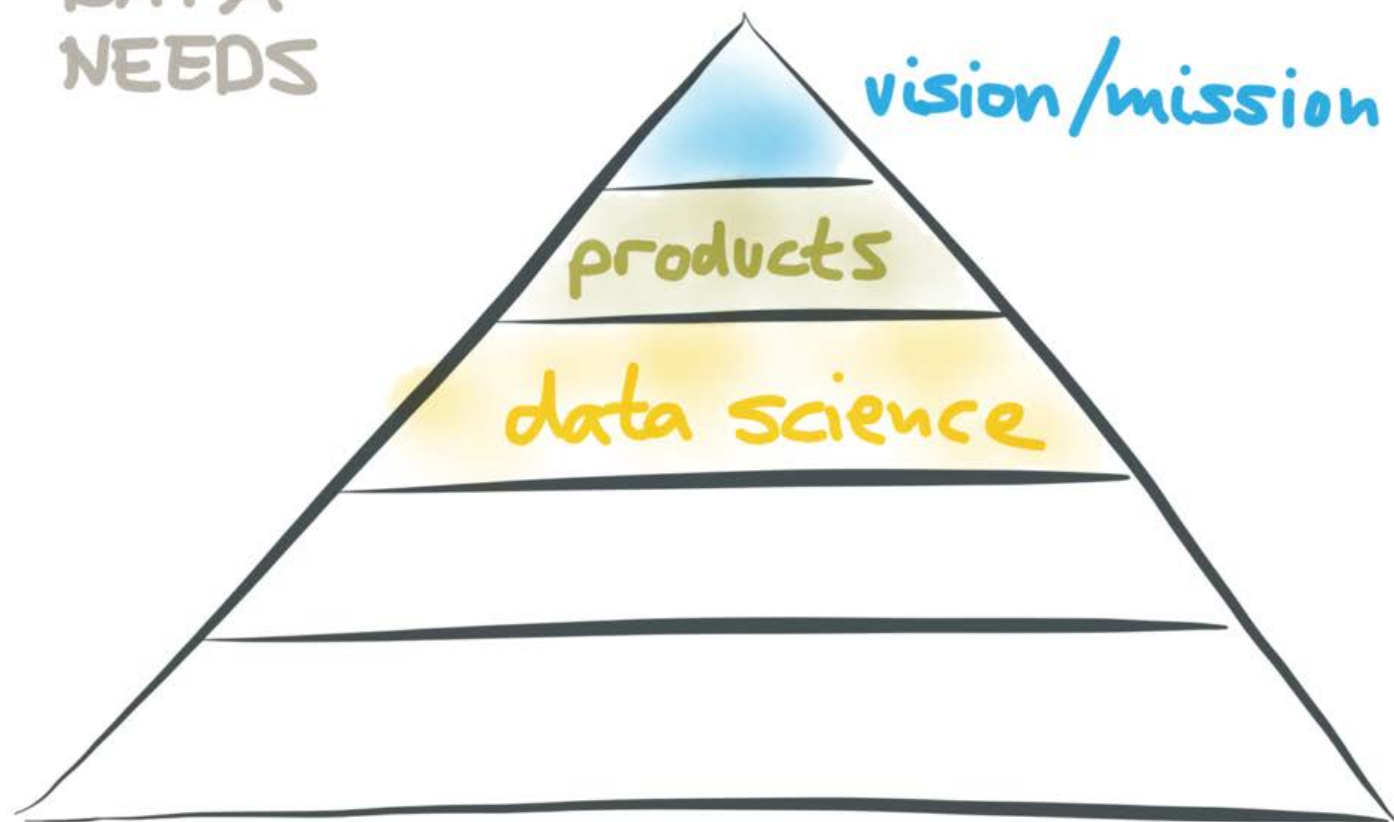
DATA  
NEEDS



DATA  
NEEDS

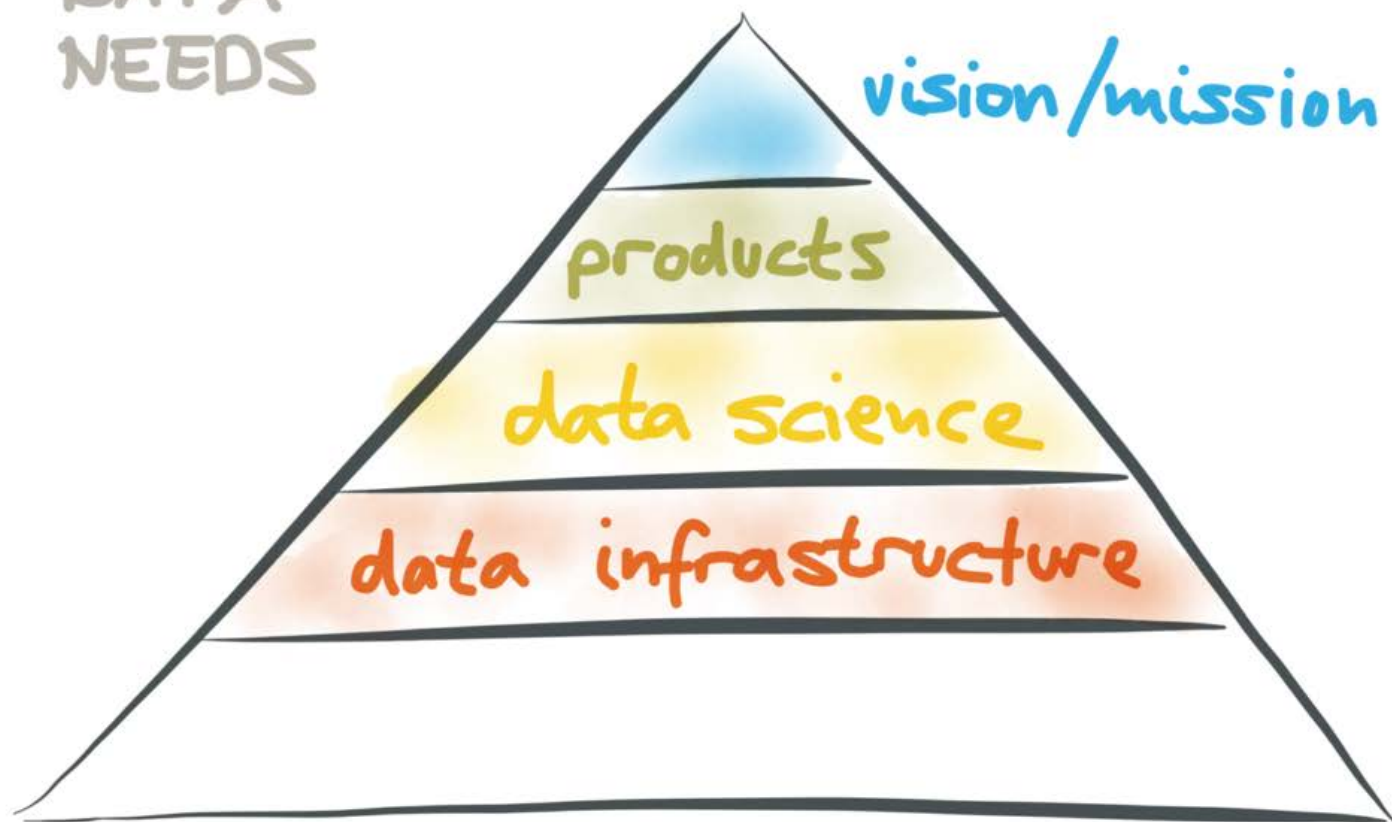


DATA  
NEEDS






DATA  
NEEDS



DATA  
NEEDS



# DATA FRAGMENTATION

Relational DBs


NoSQL DBs

Log files

Message queue

Search indexes


Monitoring



**E**XTRACT

**T**RANSFORM

**L**OAD



1990

Data warehousing

1990

Data warehousing



2008

Hadoop & MapReduce

1990

Data warehousing

- Drop relational assumption
- Programmability
- Open Source

2008

Hadoop & MapReduce



1990

Data warehousing

- Drop relational assumption
- Programmability
- Open source

2008

Hadoop & MapReduce

2015

Kafka & streaming data

1990

Data warehousing

- Drop relational assumption
- Programmability
- Open source


2008

Hadoop & MapReduce

- Batch  $\Rightarrow$  real-time
- Daily  $\Rightarrow$  continuous

2015

Kafka & streaming data



Kafka



APACHE  
SOFTWARE  
FOUNDATION

Event  
streams



"something happened"

Event

streams


subscribe to it

216.58.210.78 - - [27/Feb/2015:17:55:11 +0000] "GET /css/typography.css HTTP/1.1" 200 3377 "http://martin.kleppmann.com/" "Mozilla/5.0 (Macintosh; Intel Mac OS X 10\_9\_5) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/40.0.2214.115 Safari/537.36"

# stream

← oldest events

most recent events →




# stream

## new events added here

← oldest events

most recent events →






stream

new events added here

← oldest events

most recent events →



real-time  
consumer  
position

(close to  
head of  
stream)

"SOMETHING HAPPENED"

- User x clicked link y  
(activity event)

# "SOMETHING HAPPENED"

- User x clicked link y  
(activity event)
- Sensor x sent reading y  
(time series tick)

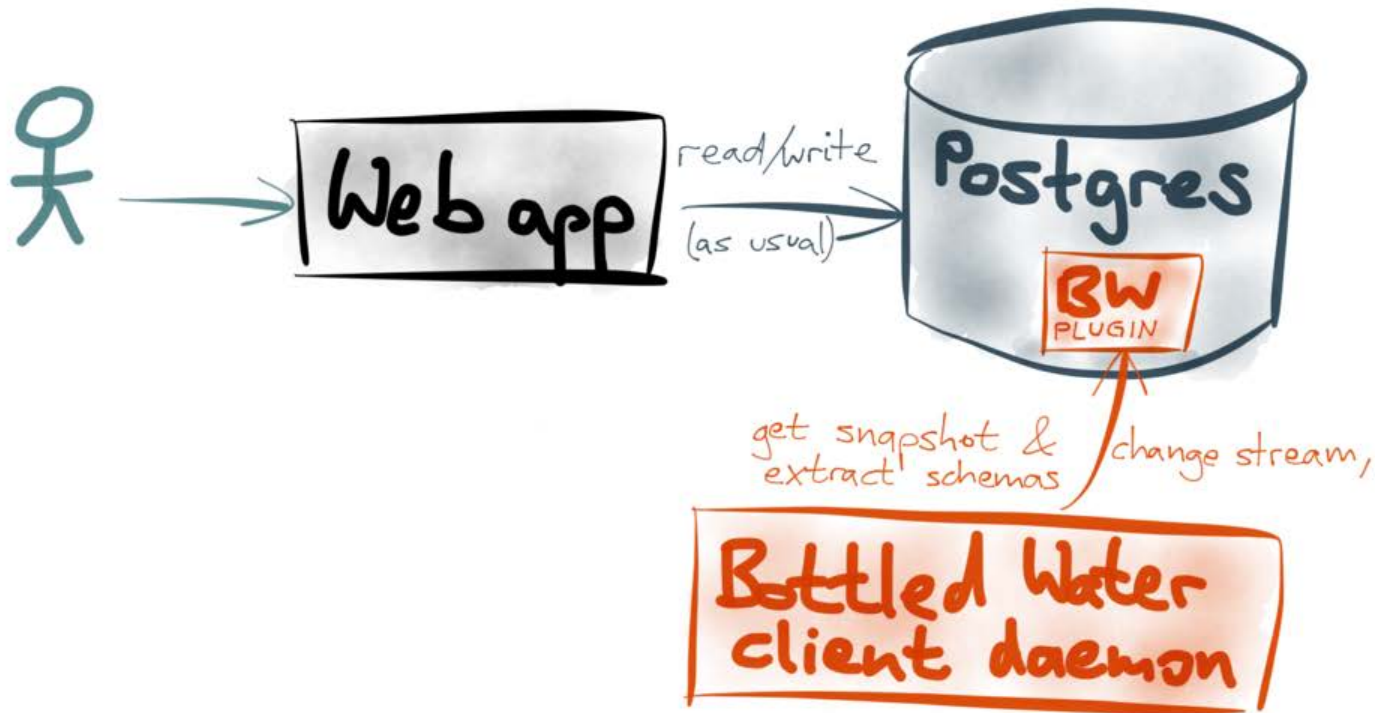
# "SOMETHING HAPPENED"

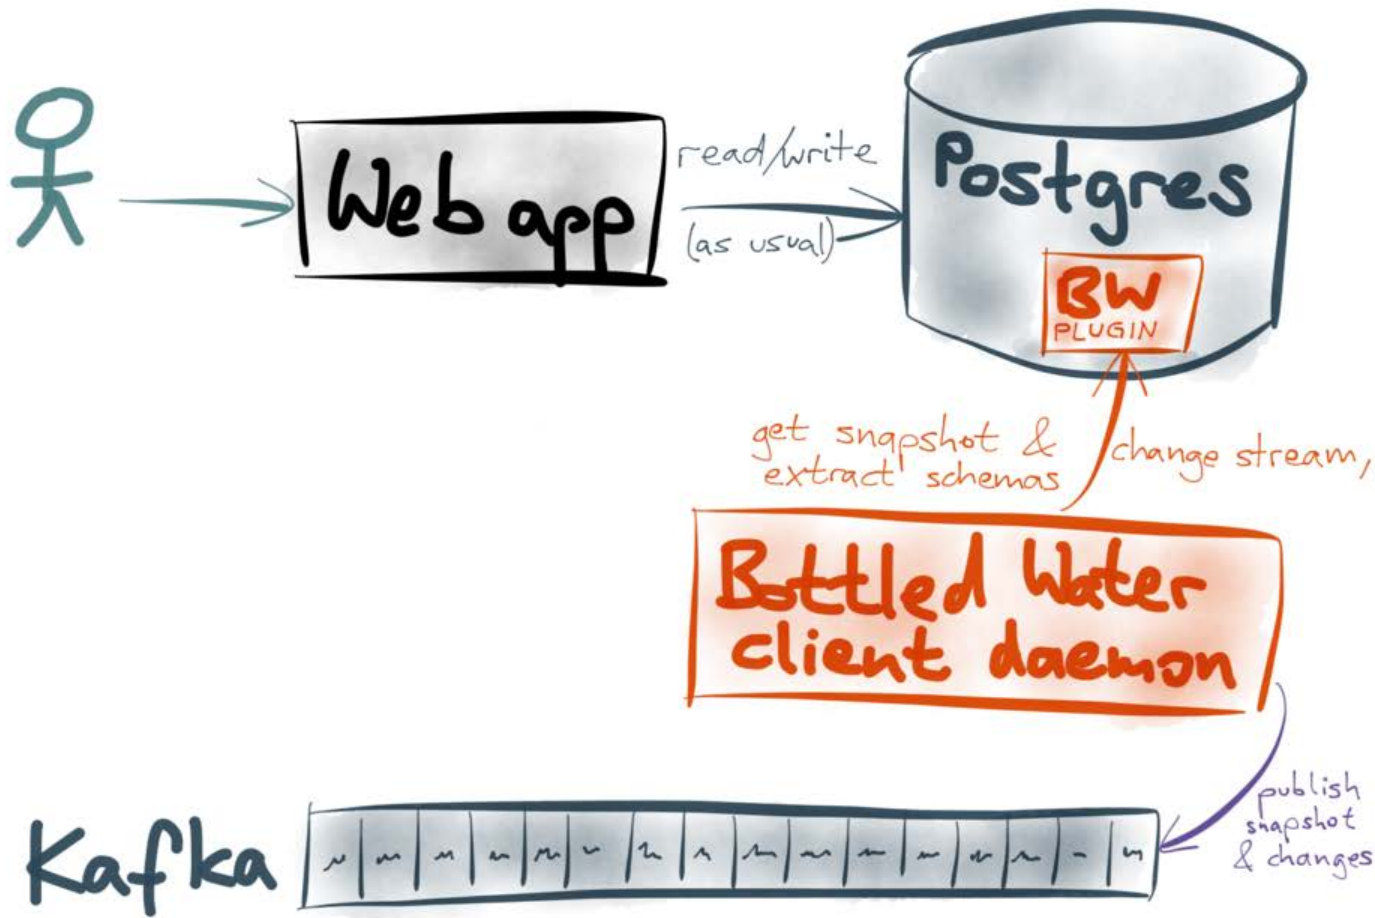
- User x clicked link y  
(activity event)
- Sensor x sent reading y  
(time series tick)
- Database record x updated to y  
(data change event)

Immutable, point in time

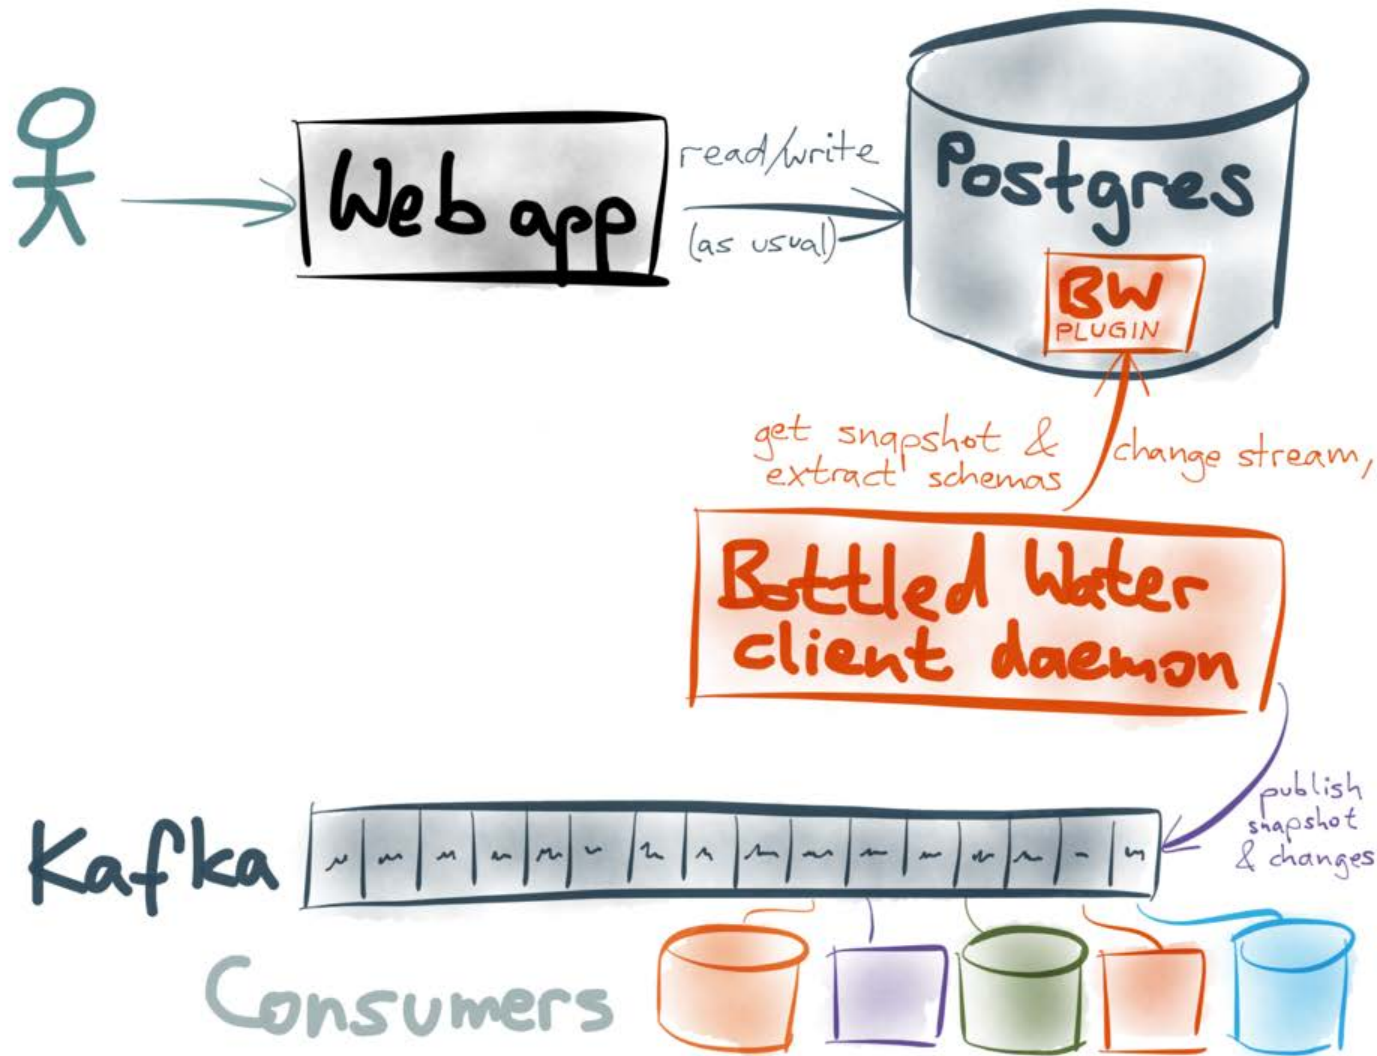


**DATA STREAMS – FRESHLY  
BOTTLED AT SOURCE**

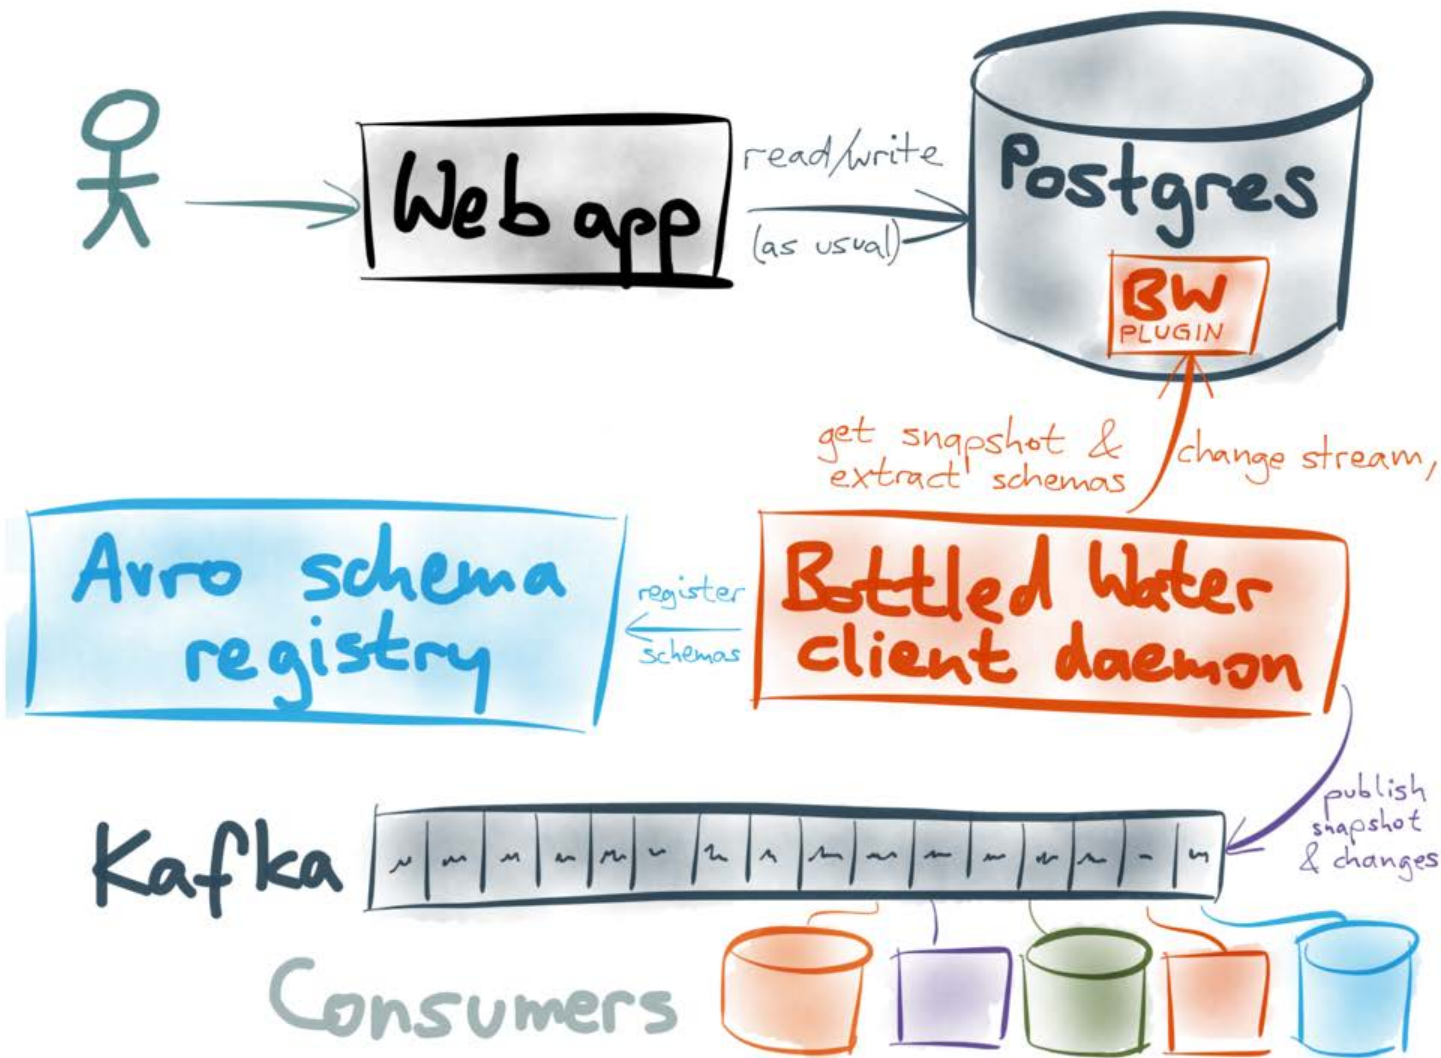























DATA  
NEEDS







→ Neha Narkhede's talk, Thu 1.15pm



1.

## Data access

Make all data available as  
Streams, including DBs



1.

## Data access

Make all data available as Streams, including DBs

2.

## Common data format

Metadata: Schema, semantics, provenance, evolution

Local optimum

Using your favorite data format

Local optimum

Using your favorite data format

Global optimum

Standardizing on one format

# JSON

+ widely supported

+ human-readable (ish)

# JSON

+ widely supported

+ human-readable (ish)

- integer / floating-point mess

# JSON

+ widely supported

+ human-readable (ish)

- integer / floating-point mess

- no binary strings

# JSON

+ widely supported

+ human-readable (ish)

- integer / floating-point mess

- no binary strings

- verbose & slow







```
record PageViewEvent {  
    long timestamp;  
    string pageURL;  
    union { IPv4Addr, IPv6Addr }  
        clientIP;  
}
```

```
record PageViewEvent {  
    /** milliseconds since epoch */  
    Long timestamp;  
    /** path and query params */  
    String pageURL;  
  
    /** IP address of client */  
    union { IPv4Addr, IPv6Addr }  
        clientIP;  
    ...  
}
```

```
record PageViewEvent {  
    long timestamp;  
    string pageURL;  
    union { IPv4Addr, IPv6Addr }  
        clientIP;  
    union { null, string }  
        sessionID = null;  
}
```

# Schema registry

PageViewEvent

v. 1

```
record PageViewEvent {  
  /** milliseconds since epoch */  
  long timestamp;  
  /** path and query params */  
  string pageURL;  
  /** client IP address */  
  union {IPv4Addr, IPv6Addr} clientIP;  
}
```

# Schema registry

PageViewEvent

v.1

v.2


```
record PageViewEvent {  
  /** milliseconds since epoch */  
  long timestamp;  
  /** path and query params */  
  string pageURL;  
  /** client IP address */  
  union {IPv4Addr, IPv6Addr} clientIP;  
}
```

```
record PageViewEvent {  
  /** milliseconds since epoch */  
  long timestamp;  
  /** path and query params */  
  string pageURL;  
  /** client IP address */  
  union {IPv4Addr, IPv6Addr} clientIP;  
  /** browser session ID from cookie */  
  union {null, string} sessionID = null;  
}
```

# Schema registry





```
record PageViewEvent {  
  /** milliseconds since epoch */  
  long timestamp;  
  /** path and query params */  
  string pageURL;  
  /** client IP address */  
  union {IPv4Addr, IPv6Addr} clientIP;  
}
```



## Message encoding












# Backward compatibility



# Forward compatibility



# Mixed versions



# Schema registry

Version number assignment

Compatibility check

Documentation

1.

Data access

Make all data available as  
Streams, including DBs

2.

Common data format


Metadata: Schema, semantics,  
provenance, evolution


1. Data access  
Make all data available as  
Streams, including DBs

2. Common data format  
Metadata: Schema, semantics,  
provenance, evolution

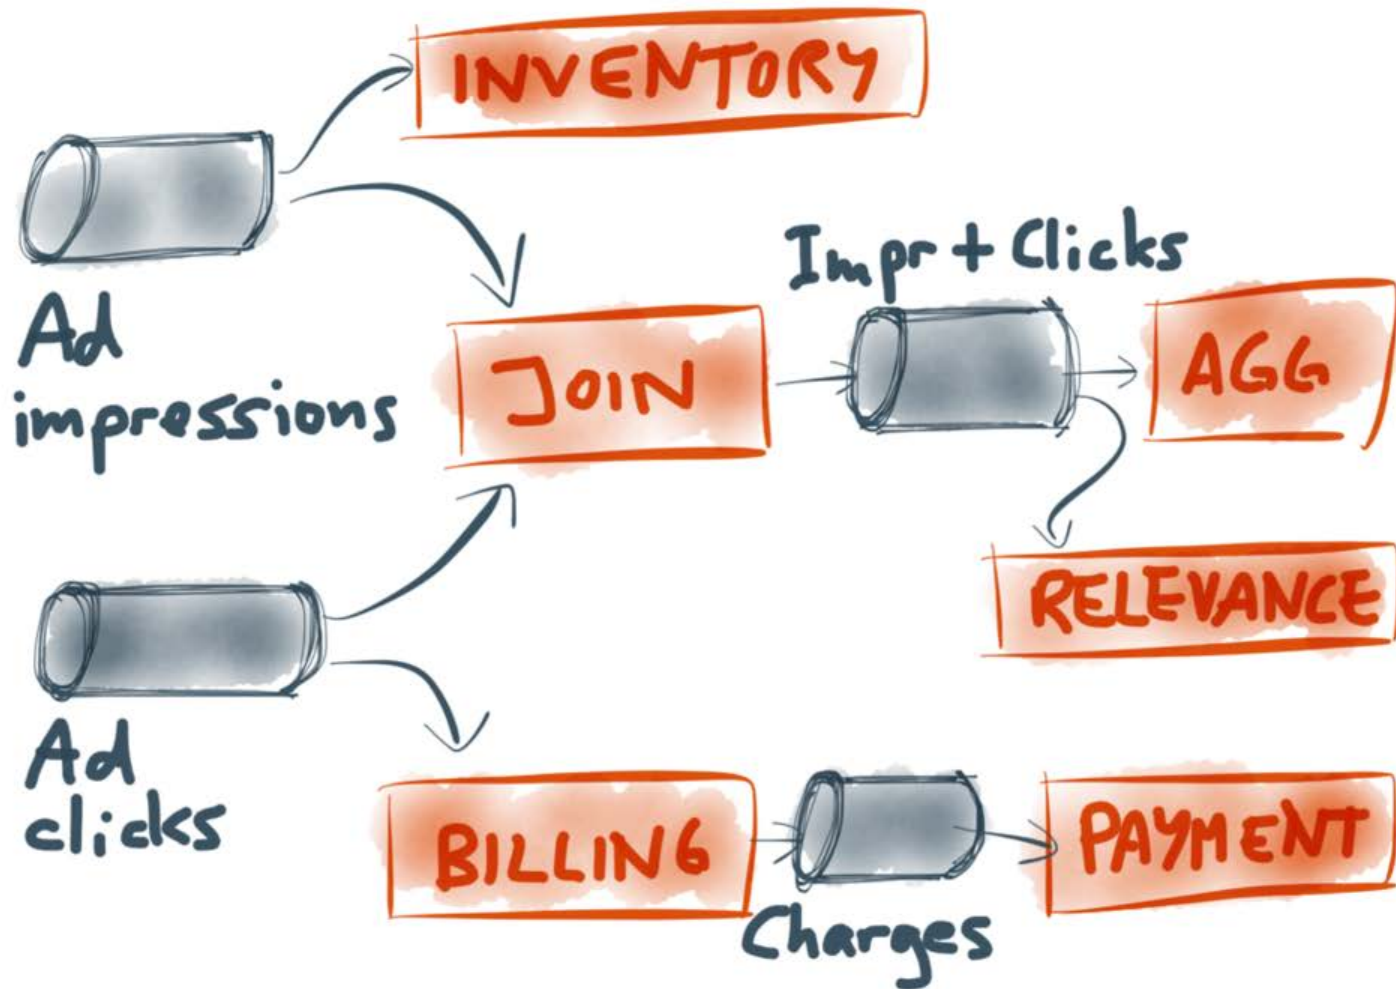
3. Composability  
Loosely coupled stream  
processors, Kafka as pipe

















Central hub  
for event  
streams









# References

1. Jay Kreps: “Putting Apache Kafka to use: A practical guide to building a stream data platform (part I).” 25 February 2015. <http://blog.confluent.io/2015/02/25/stream-data-platform-1/>
2. Gwen Shapira: “The problem of managing schemas,” 4 November 2014. <http://radar.oreilly.com/2014/11/the-problem-of-managing-schemas.html>
3. Martin Kleppmann: “Schema evolution in Avro, Protocol Buffers and Thrift,” 5 December 2012. <http://martin.kleppmann.com/2012/12/05/schema-evolution-in-avro-protocol-buffers-thrift.html>
4. Martin Kleppmann: “Bottled Water: Real-time integration of PostgreSQL and Kafka.” 23 April 2015. <http://blog.confluent.io/2015/04/23/bottled-water-real-time-integration-of-postgresql-and-kafka/>
5. Martin Kleppmann: “Designing data-intensive applications.” O’Reilly Media, to appear. <http://dataintensive.net>
6. Shirshanka Das, Chavdar Botev, Kapil Surlaker, et al.: “All Aboard the Databus!,” at *ACM Symposium on Cloud Computing (SoCC)*, October 2012. <http://www.socc2012.org/slides-das.pdf>

office  
hours



Office hours:


**5.25pm today**

O'Reilly Booth  
Expo Hall



Discount code: **TS2015**  
50% off ebooks

 **dataintensive.net**



Free signed  
copies!

Thu 3.35pm  
@O'Reilly Booth  
Expo Hall